

# Homework 2: Convolutional Neural Network

## Code:

#Author - Het Pathak

```
import argparse
import os
import random
import time
import matplotlib.pyplot as plt
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from thop import profile
import cv2

# Check if CUDA is available
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Set random seed for reproducibility
def set_seed(seed):
    random.seed(seed)
    os.environ['PYTHONHASHSEED'] = str(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    if torch.cuda.is_available():
        torch.cuda.manual_seed(seed)
        torch.backends.cudnn.deterministic = True
        torch.backends.cudnn.benchmark = False

# Define the CNN model
class CNNModel(nn.Module):
    def __init__(self):
        super(CNNModel, self).__init__()
        # First CONV layer with specific requirements:
        # filter size: 5x5, stride: 1, no padding, 32 filters
        self.conv1 = nn.Conv2d(3, 32, kernel_size=5, stride=1, padding=0)

        # Additional CONV layers
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, stride=1, padding=1)
```

```

# Batch normalization layers
self.bn1 = nn.BatchNorm2d(32)
self.bn2 = nn.BatchNorm2d(64)
self.bn3 = nn.BatchNorm2d(128)

# Pooling layer
self.pool = nn.MaxPool2d(2, 2)

# Dropout for regularization
self.dropout = nn.Dropout(0.25)

# Calculate input size for first FC layer
# Input: 32x32, after conv1: 28x28, after pool: 14x14
# After conv2: 14x14, after pool: 7x7
# After conv3: 7x7, after pool: 3x3
# So 128 channels with 3x3 feature maps = 128*3*3 = 1152
self.fc1 = nn.Linear(128 * 3 * 3, 512)
self.fc2 = nn.Linear(512, 128)
self.fc3 = nn.Linear(128, 10)

def forward(self, x):
    # First convolutional layer with ReLU and batch norm
    x = self.pool(F.relu(self.bn1(self.conv1(x))))

    # Second convolutional layer with ReLU and batch norm
    x = self.pool(F.relu(self.bn2(self.conv2(x))))

    # Third convolutional layer with ReLU and batch norm
    x = self.pool(F.relu(self.bn3(self.conv3(x))))

    # Flatten the output for the fully connected layers
    x = x.view(-1, 128 * 3 * 3)

    # Fully connected layers with ReLU and dropout
    x = F.relu(self.fc1(x))
    x = self.dropout(x)
    x = F.relu(self.fc2(x))
    x = self.dropout(x)
    x = self.fc3(x)

    return x

def get_conv1_output(self, x):
    # Get the output of the first convolutional layer
    x = self.conv1(x)
    return x

# Load the CIFAR-10 dataset
def load_data():
    # Define data transformations
    transform_train = transforms.Compose([

```

```

        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
    ])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

# Load training and test datasets
trainset = torchvision.datasets.CIFAR10(
    root='./data', train=True, download=True, transform=transform_train)
trainloader = torch.utils.data.DataLoader(
    trainset, batch_size=128, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(
    root='./data', train=False, download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(
    testset, batch_size=100, shuffle=False, num_workers=2)

# CIFAR-10 classes
classes = ('airplane', 'car', 'bird', 'cat', 'deer',
           'dog', 'frog', 'horse', 'ship', 'truck')

return trainloader, testloader, classes

# Test the model on the full test set
def test_model(model, testloader, criterion=None):
    model.eval()
    correct = 0
    total = 0
    running_loss = 0.0
    with torch.no_grad():
        for data in testloader:
            images, labels = data
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()

    # Calculate test loss
    if criterion is not None:
        loss = criterion(outputs, labels)
        running_loss += loss.item()

    accuracy = 100 * correct / total
    test_loss = running_loss / len(testloader) if criterion is not None else 0
    return accuracy, test_loss

```

```

def train(args=None):
    if not os.path.exists('./model'):
        os.makedirs('./model')

    seeds = [42, 123, 1000]
    accuracies = []
    best_accuracy = 0
    best_model_state = None

    for seed_idx, seed in enumerate(seeds):
        set_seed(seed)
        trainloader, testloader, classes = load_data()
        model = CNNModel().to(device)
        criterion = nn.CrossEntropyLoss()
        optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9, weight_decay=5e-4)
        scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=200)

        num_epochs = 10
        train_accuracies = [] # List to store training accuracies
        test_accuracies = [] # List to store testing accuracies
        print(f"\nSeed: {seed}")
        print("Epoch TrainLoss TrainAccuracy TestAccuracy TestLoss")

        for epoch in range(num_epochs):
            model.train()
            running_loss = 0.0
            correct = 0
            total = 0

            for i, data in enumerate(trainloader, 0):
                inputs, labels = data
                inputs, labels = inputs.to(device), labels.to(device)
                optimizer.zero_grad()
                outputs = model(inputs)
                loss = criterion(outputs, labels)
                loss.backward()
                optimizer.step()

                _, predicted = torch.max(outputs.data, 1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()
                running_loss += loss.item()

            train_accuracy = 100 * correct / total
            test_accuracy, test_loss = test_model(model, testloader, criterion)
            model.train()

            # Store accuracies for plotting
            train_accuracies.append(train_accuracy)
            test_accuracies.append(test_accuracy)

```

```

    avg_loss = running_loss / len(trainloader)
    print(
        f'{epoch + 1:<6} {avg_loss:.3f}    {train_accuracy:.2f}%    {test_accuracy:.2f}%
{test_loss:.3f}')"

    # Track best model
    if test_accuracy > best_accuracy:
        best_accuracy = test_accuracy
        best_model_state = model.state_dict()

    scheduler.step()

# Accuracy plot (only for the first seed)
if seed_idx == 0:
    plt.figure(figsize=(10, 6))
    epochs = range(1, num_epochs + 1)
    plt.plot(epochs, train_accuracies, 'b-o', label='Training Accuracy')
    plt.plot(epochs, test_accuracies, 'r-s', label='Testing Accuracy')
    plt.title('Training and Testing Accuracy vs. Epochs')
    plt.xlabel('Epochs')
    plt.ylabel('Accuracy (%)')
    plt.grid(True)
    plt.legend()
    plt.savefig('./model/accuracy_plot.png')
    print(f"Training accuracy plot saved to './model/accuracy_plot.png'")

# Save the final test accuracy
final_accuracy = test_accuracies[-1]
accuracies.append(final_accuracy)

# Save only the best model
torch.save(best_model_state, './model/cifar_model.pt')
print(f"\nModel saved in file: ./model/cifar_model.pt")

print("\n--- Training Summary ---")
for i, (seed, accuracy) in enumerate(zip(seeds, accuracies)):
    print(f"Seed {seed}: Test Accuracy = {accuracy:.2f}%")

mean_accuracy = np.mean(accuracies)
std_accuracy = np.std(accuracies)
print(f"Mean Accuracy: {mean_accuracy:.2f}%")
print(f"Standard Deviation: {std_accuracy:.2f}%")

# Test function to predict a single image
def test(image_path):
    # Load the saved model
    model = CNNModel().to(device)
    model.load_state_dict(torch.load('./model/cifar_model.pt'))
    model.eval()

```

```

# CIFAR-10 classes
classes = ('airplane', 'car', 'bird', 'cat', 'deer',
           'dog', 'frog', 'horse', 'ship', 'truck')

# Load and preprocess the image
image = cv2.imread(image_path)
if image is None:
    print(f"Error: Could not read image file: {image_path}")
    return

# Resize to 32x32 and convert BGR to RGB
image = cv2.resize(image, (32, 32))
image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Convert to tensor and normalize
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

image_tensor = transform(image).unsqueeze(0).to(device)

# Predict
with torch.no_grad():
    output = model(image_tensor)
    _, predicted = torch.max(output, 1)
    pred_class = classes[predicted.item()]

print(f"Predicted class: {pred_class}")

# Get the output of the first convolutional layer
conv_output = model.get_conv1_output(image_tensor)
conv_output = conv_output.detach().cpu().numpy()[0]

# Visualize the output of each filter in the first CONV layer
plt.figure(figsize=(8, 8))
for i in range(32):
    plt.subplot(4, 8, i + 1)
    plt.imshow(conv_output[i], cmap='viridis')
    plt.axis('off')

plt.tight_layout()
plt.savefig('CONV_rslt.png')

# Function to load and test ResNet-20
def test_resnet20():
    # Import ResNet-20 model
    from resnet20_cifar import resnet20

    # Load pretrained model
    model = resnet20().to(device)

```

```

model_path = "./model/resnet20_cifar10_pretrained.pt"
model.load_state_dict(torch.load(model_path, map_location=device))
model.eval()

# Load test data with the same normalization as used in training
transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

testset = torchvision.datasets.CIFAR10(
    root='./data', train=False, download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(
    testset, batch_size=100, shuffle=False, num_workers=2)

# Test the model
accuracy, test_loss = test_model(model, testloader)
print(f"ResNet-20 Test Accuracy: {accuracy:.2f}%")

# Count parameters and MACs using THOP
def count_parameters_and_macs(model, model_name):
    # Create a dummy input
    dummy_input = torch.randn(1, 3, 32, 32).to(device)

    # Count MACs and parameters
    macs, params = profile(model, inputs=(dummy_input,))

    # print(f"{model_name} - Parameters: {params / 1e6:.2f}M, MACs: {macs / 1e6:.2f}M")
    print(f"{model_name} - Parameters: {int(params)}, MACs: {int(macs)}")

# Test inference speed
def inference_speed_test(args=None):
    # Load both models
    # 1. Custom CNN model
    cnn_model = CNNModel().to(device)
    cnn_model.load_state_dict(torch.load('./model/cifar_model.pt'))
    cnn_model.eval()

    # 2. ResNet-20 model
    from resnet20_cifar import resnet20
    resnet_model = resnet20().to(device)
    model_path = "./model/resnet20_cifar10_pretrained.pt"
    resnet_model.load_state_dict(torch.load(model_path, map_location=device))
    resnet_model.eval()

    # Create a dummy input
    dummy_input = torch.randn(1, 3, 32, 32).to(device)

    # Warm-up iterations
    for _ in range(10):

```

```

_ = cnn_model(dummy_input)
_ = resnet_model(dummy_input)

# Test CNN model inference speed
torch.cuda.synchronize() if torch.cuda.is_available() else None
start_time = time.time()
num_iterations = 1000

with torch.no_grad():
    for _ in range(num_iterations):
        _ = cnn_model(dummy_input)

torch.cuda.synchronize() if torch.cuda.is_available() else None
end_time = time.time()
cnn_inference_time = (end_time - start_time) / num_iterations * 1000 # in milliseconds

# Test ResNet-20 model inference speed
torch.cuda.synchronize() if torch.cuda.is_available() else None
start_time = time.time()

with torch.no_grad():
    for _ in range(num_iterations):
        _ = resnet_model(dummy_input)

torch.cuda.synchronize() if torch.cuda.is_available() else None
end_time = time.time()
resnet_inference_time = (end_time - start_time) / num_iterations * 1000 # in milliseconds

# Print results
print(f"\nInference Speed Test Results (average over {num_iterations} iterations):")
print(f"Custom CNN model: {cnn_inference_time:.4f} ms per inference")
print(f"ResNet-20 model: {resnet_inference_time:.4f} ms per inference")

# Also count parameters and MACs for both models
print("\nModel Complexity Analysis:")
count_parameters_and_macs(cnn_model, "Custom CNN")
count_parameters_and_macs(resnet_model, "ResNet-20")

# Main function to parse arguments
def main():
    parser = argparse.ArgumentParser(description='CNN Classifier for CIFAR-10')
    parser.add_argument('command', type=str, help='train, test, or resnet20')
    parser.add_argument('image_path', nargs='?', help='Path to the image for testing')

    args = parser.parse_args()

    if args.command == 'train':
        train()
    elif args.command == 'test':
        if args.image_path is None:
            print("Error: Please provide an image path for testing")

```



```

        return
        test(args.image_path)
    elif args.command == 'resnet20':
        test_resnet20()
    elif args.command == 'inference':
        inference_speed_test()

if __name__ == '__main__':
    main()

```

## Network Description:

### ➤ Three Convolutional Layers

- conv1: 32 filters,  $5 \times 5$  kernel, stride=1, no padding.
- conv2: 64 filters,  $3 \times 3$  kernel, stride=1, padding=1.
- conv3: 128 filters,  $3 \times 3$  kernel, stride=1, padding=1.

### ➤ Batch Normalization (bn1, bn2, bn3)

- Normalizes activations to speed up training and stabilize learning.

### ➤ Max Pooling (pool)

- **$2 \times 2$  pooling after each convolutional layer** to reduce feature map size.

### ➤ Dropout Regularization (0.25 probability)

- Applied before fully connected layers to prevent overfitting.

### ➤ Fully Connected Layers (FC Layers)

- fc1: **512 neurons**
- fc2: **128 neurons**
- fc3: **10 neurons** (corresponding to CIFAR-10 classes).

### ➤ Activation Function:

- Uses **ReLU** activation for non-linearity in all layers except the last one.
- The last layer (fc3) outputs raw logits (before softmax).

# Training Output:

```
Terminal Local (2) x + v
(.venv) PS C:\Users\hetpa\PycharmProjects\DeepLearningProjs> python CNNClassify.py train

Seed: 42
Epoch TrainLoss TrainAccuracy TestAccuracy TestLoss
1 1.597 40.65% 54.29% 1.330
2 1.230 55.56% 59.98% 1.130
3 1.086 61.36% 66.14% 0.960
4 0.985 65.16% 68.84% 0.887
5 0.917 67.80% 71.43% 0.814
6 0.859 70.05% 73.46% 0.751
7 0.811 71.80% 74.50% 0.715
8 0.774 72.93% 75.11% 0.704
9 0.746 74.10% 73.53% 0.774
10 0.715 75.21% 77.43% 0.653
11 0.696 75.74% 78.14% 0.629
12 0.667 76.96% 77.63% 0.640
13 0.647 77.49% 77.25% 0.652
14 0.630 78.27% 79.75% 0.583
15 0.617 78.79% 81.05% 0.563
16 0.604 79.18% 80.40% 0.568
17 0.588 79.86% 80.27% 0.584
18 0.573 80.18% 81.07% 0.554
19 0.558 80.95% 79.74% 0.592
20 0.557 81.01% 78.41% 0.640
Training accuracy plot saved to './model/accuracy_plot.png'
```

```
Terminal Local (2) x + v
Seed: 123
Epoch TrainLoss TrainAccuracy TestAccuracy TestLoss
1 1.629 39.20% 53.01% 1.268
2 1.248 54.93% 58.53% 1.130
3 1.091 60.94% 62.88% 1.057
4 0.990 64.68% 70.21% 0.858
5 0.917 67.49% 71.06% 0.833
6 0.869 69.51% 72.20% 0.782
7 0.817 71.56% 72.18% 0.797
8 0.777 72.85% 74.26% 0.749
9 0.741 74.20% 77.30% 0.651
10 0.717 74.90% 77.05% 0.671
11 0.688 75.99% 75.17% 0.709
12 0.669 76.83% 76.45% 0.688
13 0.650 77.51% 79.16% 0.597
14 0.635 77.88% 80.66% 0.558
15 0.615 78.76% 80.37% 0.579
16 0.598 79.39% 80.98% 0.545
17 0.595 79.51% 80.71% 0.564
18 0.570 80.34% 81.89% 0.529
19 0.563 80.49% 81.33% 0.525
20 0.545 81.10% 81.69% 0.530
```

```
Terminal Local (2) x + v
Seed: 1000
Epoch TrainLoss TrainAccuracy TestAccuracy TestLoss
1 1.628 39.04% 53.42% 1.297
2 1.249 54.79% 61.30% 1.091
3 1.098 60.90% 67.25% 0.927
4 0.981 65.28% 69.82% 0.869
5 0.903 68.45% 72.57% 0.784
6 0.854 70.00% 73.87% 0.748
7 0.801 72.04% 75.73% 0.701
8 0.765 73.14% 76.36% 0.698
9 0.733 74.60% 75.89% 0.705
10 0.704 75.63% 77.97% 0.637
11 0.684 76.36% 77.24% 0.657
12 0.658 77.26% 79.05% 0.619
13 0.633 78.10% 78.77% 0.622
14 0.619 78.66% 79.76% 0.585
15 0.604 79.25% 81.17% 0.546
16 0.587 79.72% 81.53% 0.541
17 0.573 80.36% 81.09% 0.561
18 0.563 80.49% 81.30% 0.545
19 0.553 80.93% 82.60% 0.510
20 0.540 81.44% 81.01% 0.556
Model saved in file: ./model/cifar_model.pt
```

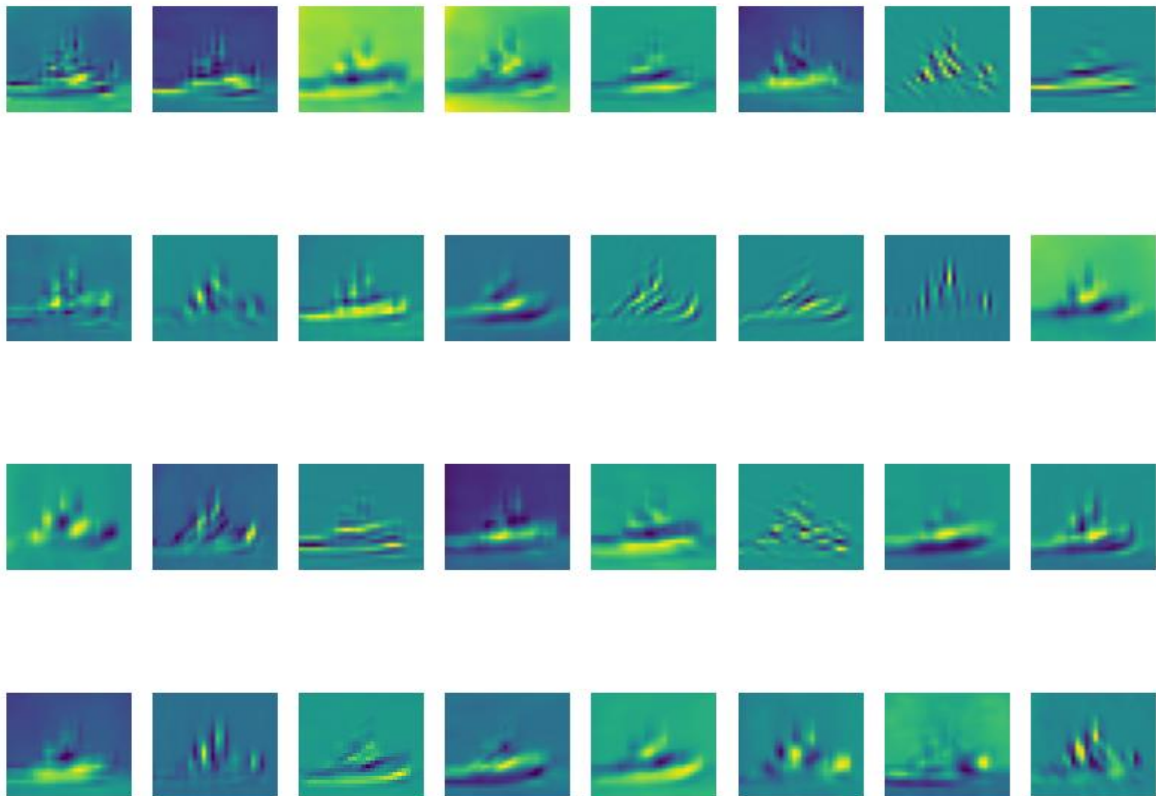
Highest Test Accuracy: 82.60%

## Testing Output:

```
Terminal  Local (2)  ×  +  ▾  
(..env) PS C:\Users\hetpa\PycharmProjects\DeepLearningProjs> python CNNclassify.py test 0008.png  
Predicted class: ship
```

- 0008.png is test image from CIFAR-10 dataset.
- random\_car.png is a random image of a car downloaded from google.

## Visualization Result- First CONV layer:




## ResNet accuracy on CIFAR-10 dataset:

```
Terminal Local (2) x + v
Predicted class: ship
(.venv) PS C:\Users\hetpa\PycharmProjects\DeepLearningProjs> python CNNclassify.py resnet20
ResNet-20 Test Accuracy: 92.05%
```

## Train/Test Accuracy vs Epochs plot:

- Plot for the 1<sup>st</sup> seed.



## Inference Speed:

```
Inference Speed Test Results (average over 1000 iterations):
Custom CNN model: 2.7776 ms per inference
ResNet-20 model: 11.2904 ms per inference
```

## Computation Costs and Number of Parameters:

```
Model Complexity Analysis:
[INFO] Register count_convNd() for <class 'torch.nn.modules.conv.Conv2d'>.
[INFO] Register count_normalization() for <class 'torch.nn.modules.batchnorm.BatchNorm2d'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.pooling.MaxPool2d'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.dropout.Dropout'>.
[INFO] Register count_linear() for <class 'torch.nn.modules.linear.Linear'>.
Custom CNN - Parameters: 752522, MACs: 9939200
[INFO] Register count_convNd() for <class 'torch.nn.modules.conv.Conv2d'>.
[INFO] Register count_normalization() for <class 'torch.nn.modules.batchnorm.BatchNorm2d'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.activation.ReLU'>.
[INFO] Register zero_ops() for <class 'torch.nn.modules.container.Sequential'>.
[INFO] Register count_avgpool() for <class 'torch.nn.modules.pooling.AvgPool2d'>.
[INFO] Register count_linear() for <class 'torch.nn.modules.linear.Linear'>.
ResNet-20 - Parameters: 272474, MACs: 41616064
```

## Training Summary:

Seed 42: Final Test Accuracy = 78.41%

Seed 123: Final Test Accuracy = 81.69%

Seed 1000: Final Test Accuracy = 81.01%

Mean Accuracy: 80.37%

Standard Deviation: 1.41%

- To check Inference speed, Computation Cost, and Number of Parameters for my CNN model and pretrained resnet20 model run: **“python CNNclassify inference”**.